

Reinforcement Learning

Aldo Faisal with contributions
by Ed Johns and Paul Bilokon

Imperial College London


July 2021 – Version 7.01

Section overview







Deep Q Learning

- 1 Motivation
- 2 Reinforcement Learning 101
- 3 Lets go Markov
- 4 Markov Decision Process
- 5 Dynamic Programming
- 8 Function Approximation
- 9 Deep Q Learning
 - Putting the Deep in Reinforcement Learning
 - Policy Gradients
 - Mathematical interlude: Gaussian Distribution
 - Back to Policy Gradients
- 10 Actor-Critic Methods


Mountain Car Video





Example: Mountain Car example





Space Invaders





Breakout





Off-Policy Function Approximation TD: The Deep RL triad

- Function approximation: A powerful, scalable way of generalising from a state space much larger than the memory and computational resources available (e.g., linear function approximation or CNNs). Leads to Deep RL.
- Bootstrapping: Update targets that include existing estimates (as in dynamic programming or TD methods) rather than relying exclusively on actual rewards and complete returns (as in MC methods). Leads to Deep TD Learning.
- Off-policy: training on a distribution of transitions other than that produced by the target policy. Sweeping through the state space and updating all states uniformly, as in dynamic programming, does not respect the target policy and is an example of off-policy training. Leads to Deep Q Learning.

Q-Learning vs Lin FA Q-Learning

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

Initialize S

Repeat (for each step of episode):

Choose A from S using policy derived from Q (e.g., ϵ -greedy)

Take action A , observe R, S'

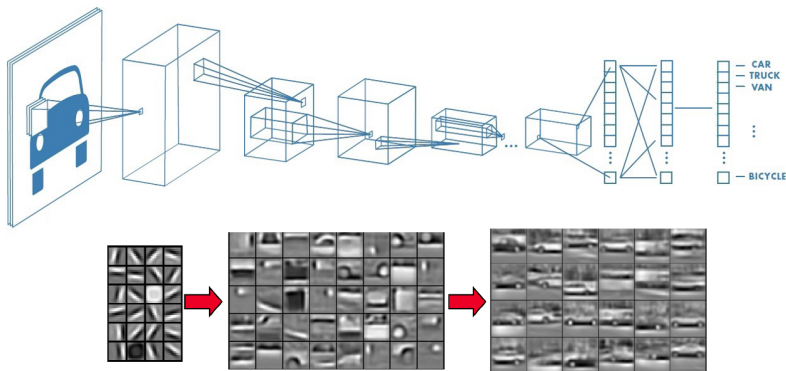
$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$;

until S is terminal

Features (programmed not learned): Deep Learning


- Neural networks (NNs) are powerful function approximators



ATTENTION: Neural Networks are normally Supervised Learners

Features (programmed not learned): Deep Learning

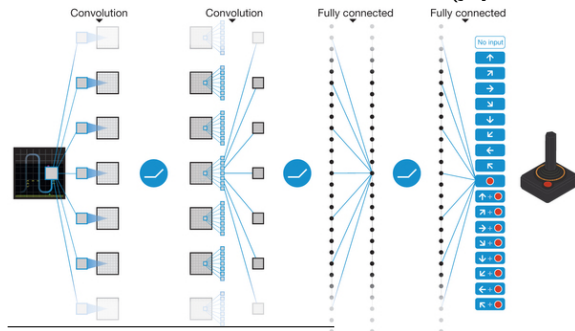
- Neural networks (NNs) are powerful function approximators
- NNs can learn features directly from (loads of) data
- Stacking layers enables learning hierarchical features (e.g. decomposing images into meaningful constituents)
- Deep learning is a subset of representation learning
- We can use it to replace hand-engineering of state space features with learning features from state data



Example: Atari DQN³

There are many function approximators, e.g.

- $Q(s, a)$ is approximated by a neural network
- Process raw pixels with convolutional layers for feature extraction
- Actions are from a discrete set (joystick commands)



³Mnih et al (2015). Human-level control through deep reinforcement learning. Nature, 518(7540), 529-533

DQN: Bringing Deep Learning into RL I

In principle it sounds easy: use convolutional neural networks to link screen shots to values.

Given our TD Q-learning update:

$$Q(s_t, a) \leftarrow Q(s_t, a) + \alpha [r_{t+1} + \gamma \max_{a_{t+1}} Q(s_{t+1}, a') - Q(s_t, a)] \quad (57)$$

we want to learn $Q(s_{t+1}, a'; w)$ as a parametrised function (a neural network) with parameters w .

We can define the TD error as our learning target that we want to reduce to zero.

$$\text{TD error}(w) = r_{t+1} + \gamma \max_{a_{t+1}} Q(s_{t+1}, a'; w) - Q(s_t, a; w) \quad (58)$$

DQN: Bringing Deep Learning into RL II

Taking the gradient of E wrt w we obtain:

$$\Delta w = \alpha [r + \gamma \max_{a_{t+1}} Q(s_{t+1}, a_{t+1}; w) - Q(s_t, a; w)] \nabla_w Q(s_t, a; w) \quad (59)$$

However, the ATARI playing problem required more engineering to solve properly

- 1 Experience Replay
- 2 Target Network
- 3 Clipping of Rewards
- 4 Skipping of Frames

DQN: Experience Replay I

The CNN is easily overfitting the latest experienced episodes and the network becomes worse at dealing with different experiences of the game world:

- ① **Complication 1: Inefficient use of interactive experience**
 - Training deep neural networks requires many updates, each has its own transition
 - But now each state transition is used only once, then discarded
 - so we would need to constantly revisit the same state transitions to train.

This is inefficient and slow.

- ② **Complication 2: Experience distribution**
 - Interactive learning produces training samples that are highly correlated (agents recent actions reflect recent policy)
 - The probability distribution of the input data changes

DQN: Experience Replay II

- The network forgets transitions that were in the not so recent past (overwriting past memory, in neuroscience known as "extinction")
- Moreover, because networks are monolithic, changes in one part of the network can have side-effects on other parts of the network [Why is this not a problem for table-based RL?]

Solution: Experience Replay⁴ (alternative use Hierarchical RL).

- Experiences (traces) are stored and replayed in mini-batches to train the neural networks on more than just the last episode.
- Instead of running Q-learning on each state/action pairs as they occur during simulation or actual experience, the experience replay buffer system stores the traces sampled.

DQN: Experience Replay III

- Mini-batches are only feasible when running multiple passes with the same data is somewhat stable with respect to the samples. I.e. the transitions should low variance/entropy for the next immediate outcomes (reward, next state) given the same state, action pair.
- Replaying past data is also a more efficient use of previous experience, by learning (training the neural network) with it multiple times. This is key when gaining real-world experience is costly (e.g. simulation time) as the Q-learning updates are incremental and do not converge quickly.

DQN: Experience Replay IV

The learning phase is thereby separated from gaining experience, and based on taking random samples from this mini-batch table. DQN interleaves the two processes - acting and learning - because improving the policy will lead to a different behaviour that we should explore actions as these are closer to optimal ones (GPI). In summary

- ➊ Reduction of correlation between experiences in updating DNN – mini-batches effectively make the samples more independently and identically distributed.
- ➋ Increased learning speed with mini-batches
- ➌ Reusing/Replaying past transitions to avoid catastrophic forgetting of associated rewards

DQN: Experience Replay V

Initialize replay memory \mathcal{D} to capacity N

Initialize action-value function Q with random weights

for episode = 1, M **do**

 Initialize state s_t

for $t = 1, T$ **do**

 With probability ϵ select a random action a_t

 otherwise select $a_t = \max_a Q^*(s_t, a; \theta)$

 Execute action a_t and observe reward r_t and state s_{t+1}

 Store transition (s_t, a_t, r_t, s_{t+1}) in \mathcal{D}

 Set $s_{t+1} = s_t$

 Sample random minibatch of transitions (s_t, a_t, r_t, s_{t+1}) from \mathcal{D}

 Set $y_j = \begin{cases} r_j & \text{for terminal } s_{t+1} \\ r_j + \gamma \max_{a'} Q(s_{t+1}, a'; \theta) & \text{for non-terminal } s_{t+1} \end{cases}$

 Perform a gradient descent step on $(y_j - Q(s_t, a_j; \theta))^2$

end for

end for

⁴Lin, L. J. (1993). Reinforcement learning for robots using neural networks (No. CMU-CS-93-103). Carnegie-Mellon Univ Pittsburgh PA School of Computer Science.

DQN: Target Network I

Whenever we run an update step on a Q-network's state we also update "near by" states (monolithic architecture + generalisation ability of CNNs).

$$Q(s_t, a) \leftarrow Q(s_t, a) + \alpha [r_{t+1} + \gamma \max_{a'} Q(s_{t+1}, a') - Q(s_t, a)] \quad (60)$$

Issue: Unstable training resulting from bootstrapping a continuous state space representation

- We may have an unstable learning process, because changes to $Q(s, a)$ change $Q(s', a)$ (with $s - s' \approx \delta$) which changes $Q(s, a)$ on the next CNN update and then in turn $Q(s', a)$ forth. This can lead to a resonance effect

DQN: Target Network II

- The effect may result in a run-away bias from bootstrapping and subsequently dominating the system numerically, causing the estimated Q values to diverge.
- In calculating the TD error calculation, the target function is changing too frequently when using convolutional neural networks.

Solution: Slow things down (as resonance dampener)

DQN: Target Network III

Using a separate target network Q' , updated every few time (e.g. every 1000 steps) with a copy of the latest learned weight parameters, controls this stability (relaxation time).

- 1 Initialise two Q networks: a main Q-network (Q), and a target network (Q')
- 2 When calculating the TD-error, use Q' , not Q
- 3 Infrequently set $Q' = Q$
- 4 This gives the highly fluctuating Q time to settle (we can measure this so called "relaxation time" empirically and set it accordingly) before updating Q'

DQN: Clipping rewards I

Clipping rewards:


- Each ATARI game has different score scales. For example, in Pong, players can get 1 point when winning the play. Otherwise, players get -1 point. However, in Space Invaders, players get 10-30 points when defeating invaders. This difference would make training unstable.
- Clipping Rewards clips the rewards, with all positive rewards are set +1 and all negative rewards are set -1.

DQN: Skipping frames I

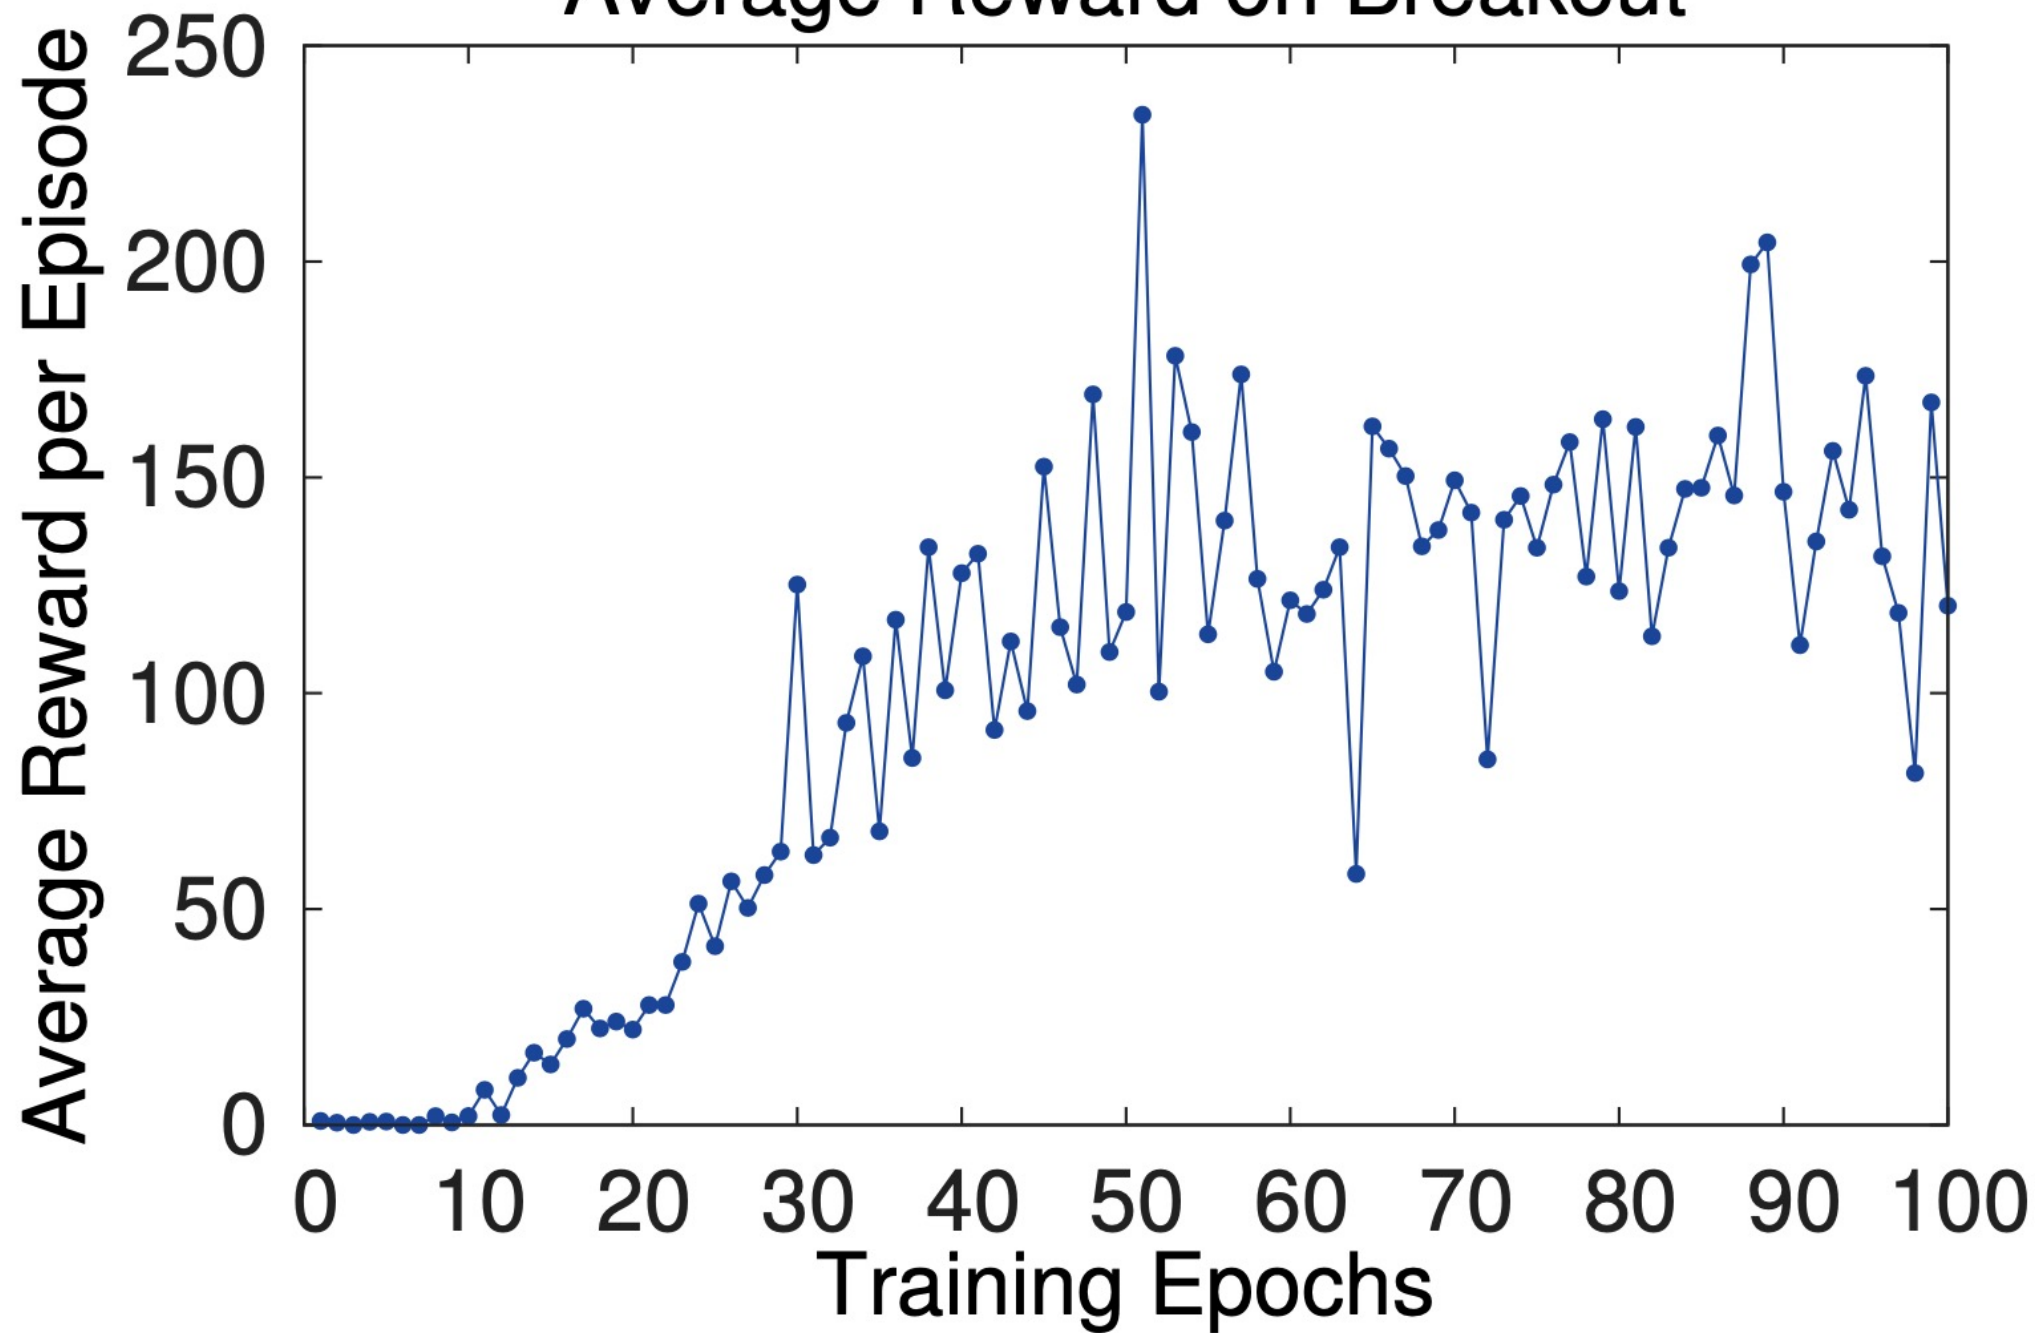
Skipping frames:

- Computer can simulate games (e.g. the ATARI simulator does is at 60 Hz) much faster than a human player would be able to react to the game, and this implies that the video game design and the required actions can be slower.
- Frame skipping uses only every 4 video game frames (i.e. 15 Hz), the past 4 frames as inputs. Reducing computational cost and accelerating training times.

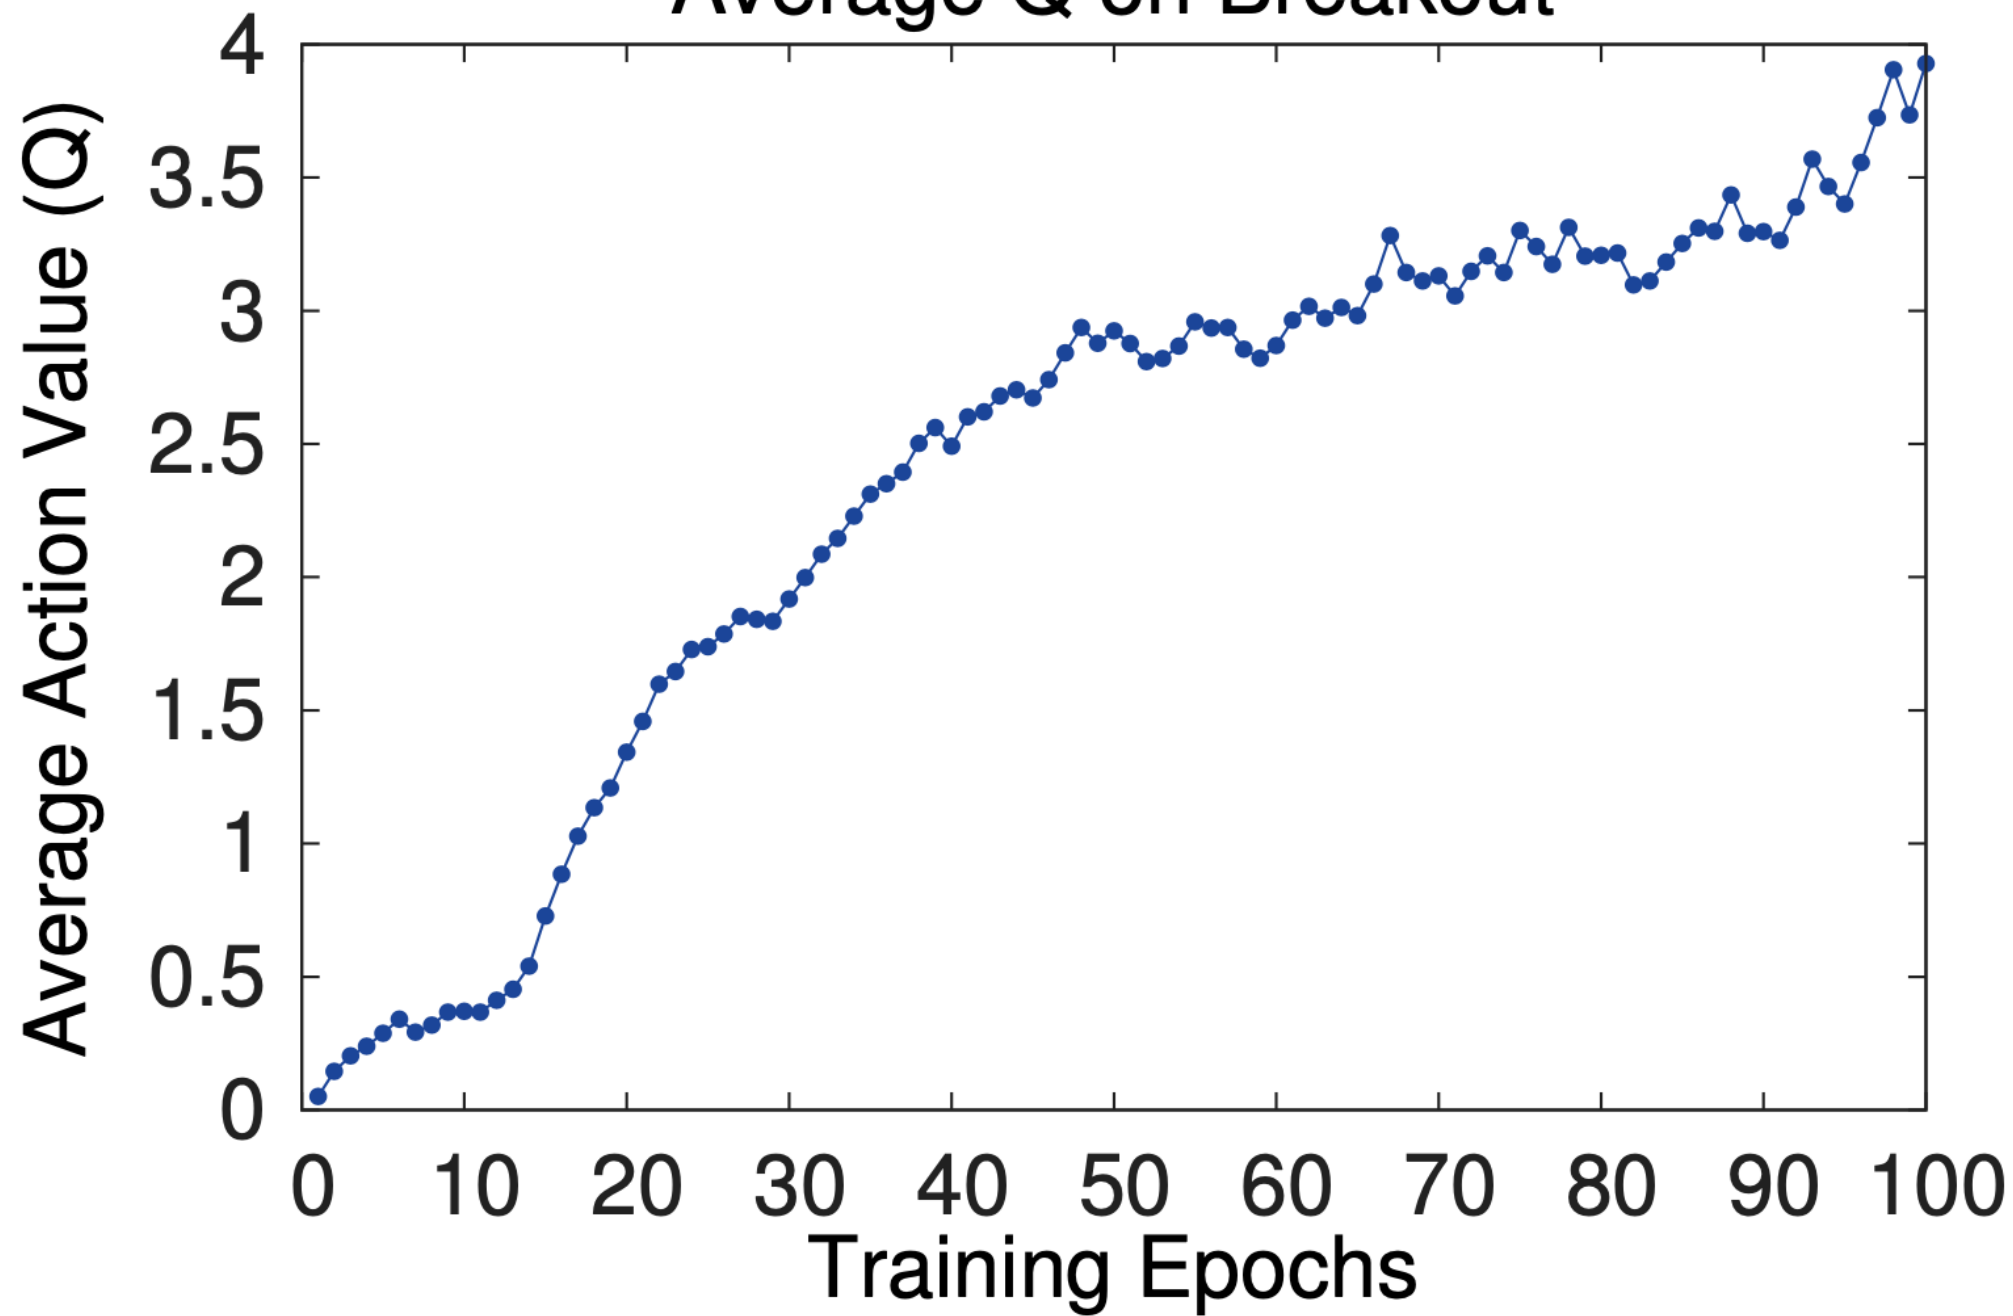
Breakout





Average Reward on Breakout



Average Q on Breakout







Algorithm 1 Deep Q-learning with Experience Replay

Initialize replay memory \mathcal{D} to capacity N

Initialize action-value function Q with random weights

for episode = 1, M **do**

 Initialise sequence $s_1 = \{x_1\}$ and preprocessed sequenced $\phi_1 = \phi(s_1)$

for $t = 1, T$ **do**

 With probability ϵ select a random action a_t

 otherwise select $a_t = \max_a Q^*(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in \mathcal{D}


 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from \mathcal{D}

 Set $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$

 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ according to equation

end for

end for



$$\nabla_{\theta_i} L_i(\theta_i) = \mathbb{E}_{s, a \sim \rho(\cdot); s' \sim \mathcal{E}} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) - Q(s, a; \theta_i) \right) \nabla_{\theta_i} Q(s, a; \theta_i) \right]$$